

AI for Software Testing and Reverse Engineering - Lab 1

INTRODUCTION

Software contains many errors that are difficult to find. Methods for locating such errors come in two flavors: static and dynamic analysis. Static analysis searches the code for patterns that indicate the presence of errors such as a use-after-free vulnerability. Dynamic analysis instead executes the code, monitors this execution, and aims to find inputs that trigger vulnerabilities such as segmentation faults. This course gives an overview of different techniques for dynamic analysis, which can be divided into three categories:

- Black-box – the code or internal software structure is not used*
- Grey-box – uses parts of the internal structure/code as guidance*
- White-box – fully interprets the internal structure/code*

*The first lab focusses on **grey-box methods**. You will start by building a simple **black-box fuzzer** that continuously probes a program with random input, hoping to cause problems. You will then improve this fuzzer by guiding it using **grey-box** details such as the branch coverage, a flow analysis, or input taints. Finally, your task is to compare your fuzzer to the state-of-the-art. The programs that we will use throughout this class come from the 2020 RERS (Rigorous Examinations of Reactive Systems) challenge. Specifically, we will be using the Reachability problems from challenge in the lab.*

- <https://rers-challenge.org/2020/index.php?page=reachProblems>*

These are highly obfuscated pieces of software that take strings as input and return strings as output. On Brightspace, you can find a folder containing a Docker file and a shellscript. The Docker file can be used to build a Docker container containing the relevant problems, as well as a multitude of programs that you will use throughout this course. Keep in mind that building the Docker container takes time and storage space of your system (estimated to take 6.7GB of space), so make sure you have enough disk space on your system before building the container.

LEARNING OUTCOMES

After completing this assignment, you will be able to:

- Build fuzzers from scratch.*
- Use instrumentation to gather grey-box information.*
- Use grey-box information to build smarter fuzzers that use search.*
- Use and understand AFL (<http://lcamtuf.coredump.cx/afl/>).*
- Analyze differences between fuzzing results.*

INSTRUCTIONS

Task 0: Preparations

Read the content on Brightspace on fuzzing, tainting, branch distance, hill-climbing and the document *docker_instruction.pdf*.

Download the folder containing the Docker file and the shell script from Brightspace. Build the container following the instructions that are given in *docker_instructions.pdf*.

The Reachability problems that we will use in the lab are located in the RERS directory. The Reachability problems are problems 11 till 19. These are highly obfuscated pieces of software that model reactive systems that read symbols ("iA", "iB", ...) from stdin and print output to stdout ("oZ", "oY", ...). Different input sequences trigger the execution of different code branches. The goal of the Reachability problems is to discover which of the error branches are reachable (which error codes can be triggered).

For the implementation of the techniques taught in class, we will work with the Java files and instrument them using JavaParser (<https://javaparser.org/>). We have already written the code instrumentation for you. The instructions for running the instrumentation can be found on the following repository:

- <https://github.com/apanichella/JavaInstrumentation>

Please **read the instructions carefully** before running the commands listed on the README and make sure **you pulled the latest changes** using "*git pull*" in the "*JavaInstrumentation*" folder. For this assignment, there is **no need** for you to do the instrumentation yourself. You only need **to write code for your solution** and use the instrumented files to test your solutions. Of course, if you are curious about how we have done the instrumentation, you can always have a look at the Java files that we have used to instrument the problem files.

A simple random fuzzer is already implemented. Compiling and executing the instrumented Java file should start this fuzzer. Your job is to make this fuzzer smarter.

Task 1: branch distance

For this task, you will have to implement the logic that can compute the sum of branch distances for a given input trace. You should **only need to modify** "*FuzzingLab.java*" for this assignment.

To save you some time and to make it easier for you to implement your solution, we have instrumented the RERS problems such that each time that we have encountered a new branch, the method *encounteredNewBranch* is called every time a *myIf* is called. The *myIf* statements are executed immediately before every if-branch.

You **will need to implement** the logic for the *encounteredNewBranch* method i.e., you would need to compute branch distances and use it to guide the fuzzer. For each branch, the condition, the value of the condition, and the line number where the branch is located

in the original file is given to the method. It is up to you to use this information in any way you see fit to make your fuzzer smarter.

The condition is represented as a (near) binary tree of *MyVar* variables. Every *MyVar* has the following attributes:

- Type of the variable (e.g. Boolean, Integer, String, Unary variable etc.).
- The primitive value of the variable (e.g. actual Boolean, numeric or string value).
- An operator symbol (e.g. =, +, - etc.) if the type of the variable is an expression.
- Two *MyVar* variables if the type of the variable is a binary expression (e.g. $a + b$).

Print some of the branches you encounter to understand how this represents an if-branch.

The “*encounteredNewBranch*” method has the following arguments:

- “*value*”: the actual value of the if-branch, so what the condition evaluates to (either true or false).
- “*line_nr*” is the line number of the if-statement in the original Java file. Hint: you can use this to compute branch coverage.

Write code that computes a branch distance for each encountered if-branch based on the “*value*” argument, i.e., the distance how far you are from flipping the branch. Run your solution against the following reachability problems for five minutes: 11-15 and 17.

Answer the following questions in **max 1 A4** and provide some details on the following:
For each problem:

- For each problem, what is the highest number of unique branches that you have visited?
- For each problem, which input was used to achieve the highest number of unique branches?

Task 2: search

Now that you are able to compute the branch distance for a given input trace, you can use this (grey-box) information to guide the fuzzer in the generation of the input traces. How to do this is entirely up to you, but we advise building a simple hill-climber:

1. Compute the sum of branch distances for a current trace
2. Try X random mutations on an input trace and compute their sum of branch distances
3. Select a mutation that lowers the sum of branch distances
4. If none exists, select a random trace
5. Update the current trace to the mutated trace with smallest sum.
6. Reset the system, execute the current trace, and iterate

Implement this simple search strategy (or try your own strategy) and try to trigger as many error codes possible for the RERS Reachability problems. Use the same set of

reachability problems that was given in Task 1. Run your “smart” fuzzer and answer the following the following questions **on max 1 A4**:

- Compared to the random fuzzer from Task 1, how many more unique branches were you able to visit using your search strategy?
- For each version of the fuzzers, compare how many error codes you were able to trigger on each of reachability problem (11-15 and 17). Which unique error codes were you able to reach using each version of the fuzzer?
- Which version managed to reach more unique error codes? For the comparisons, consider plotting the convergence graph (e.g. number of unique error codes over time) for each version of the fuzzer.

The task below is served as a preparation for the final assignment and will not be graded for the this assignment.

Task 3: AFL

For this task, you are asked to **apply AFL to the same set of reachability problems** that were given in Task 1. To keep it fair, also run AFL **for five minutes for each reachability problem**. AFL is a state-of-the-art fuzzer that uses genetic algorithms and branch coverage for guidance. Instructions on how to set up AFL for the RERS problems are available on GitHub:

- https://github.com/apanichella/javaInstrumentation/blob/main/docs/fuzzing_rers.md

Compare the performance of AFL against both the random fuzzer and your smart fuzzer. The findings/crashes directory contains all crashes AFL found, in this case there are the reachability errors. By running these through the normal (uninstrumented) code, you can obtain all the found reachability statements. To make it easier for you, we have written an script that collects the error codes that AFL managed to trigger for a given reachability problem. The script can be found here:

- https://github.com/apanichella/javaInstrumentation/blob/main/docs/fuzzing_rers.md (Specifically "Retrieving Results" section)

In your comparison answer the following questions:

- Does AFL reach more unique reachability statements (error codes) than your own smart fuzzer on the same set of problems and time constraint? What about the number of unique branches?
- Investigate the traces that were used by AFL to find the error codes. Did your own fuzzer also generate the same traces?

For the comparison, you can again consider to plot the convergence graph. Write down all your findings!

RESOURCES

The lab is run in a Docker container, the Docker container can be build using the Docker file that is provided on Brightspace. For the first lab, you will need the Reachability problems from the RERS challenge. These will be downloaded when the Docker container is being built. Copy the Reachability problems into the *JavaInstrumentation* folder to make instrumentation easier.

You should read and understand the fuzzingbook chapter on mutation-based fuzzing:
<https://www.fuzzingbook.org/html/MutationFuzzer.html>

PRODUCTS

A small report of max two A4 pages, excluding visualization, answering the questions from the 2 tasks. The report can be extended with at most two A4 pages if visualizations are used. Also provide an archive (tar/zip) containing your version of "*FuzzingLab.java*" and the code for computing the results.

Make sure you have also provided some instructions on how to run your code.

ASSESSMENT CRITERIA

You can either pass or fail this assignment. We expect everyone to pass. You have to complete all lab assignments. Submissions of which the report text does not fit into three A4 pages (excluding visualizations) will not be evaluated; the submission will automatically receive a fail.

Code that does not compile/run will not be evaluated; the submission will automatically receive a fail.

Your work will be evaluated based on its completeness (having done all tasks), the correctness of the implementation, and demonstration that you understand the results in the analysis.

When deemed insufficient, you will receive feedback and will be given a one-week grace period to fix any shortcomings.

SUPERVISION AND HELP

There will be lab sessions every Wednesday, where the teachers and TAs will be available to answer any questions you may have. The preferred way to ask questions is through Mattermost.

SUBMISSION AND FEEDBACK

The submission is through Brightspace. You will receive feedback within one week after the deadline.